

Pairwise win-rates (Fleiss' κ : humans 0.13, audio 0.39, text 0.45)
human $\leftrightarrow$ audio-LLM consensus $\kappa = 0.64$ (substantial)

