



Our pipeline vs general-audio reference captioners (MusicCaps test)

Legacy metrics

AAC metrics



Legend: GPT-2 (S1) (dark blue), GPT-2 (best dec.) (medium blue), T5 (best dec.) (light blue), Audio Flamingo † (red with diagonal stripes), Qwen2-Audio † (orange with diagonal stripes), Qwen3-Omni Captioner † (yellow with diagonal stripes).

Blue = our pipeline (comparison anchor); warm = general audio models. † hatched: may have seen MusicCaps in training (leakage upper bound, not a clean held-out score).